

**COMP281**

# **Principles of C and memory management**

**lecture 6**

**Dr. Frans Oliehoek  
Department of Computer Science  
University of Liverpool**

# Last Time

- Pointers
  - A pointer is denoted with a \*
  - \* is used to dereference a pointer
  - The address of variable can be given by &variable
- Arrays & Pointers
  - “*pointer arithmetic* and *array indexing* are equivalent in C, *pointers* and *arrays* are different.”
- Start Heap Memory allocation
  - malloc, free


# Today

- Wrap up malloc/free
- Debugging heap memory allocation: `valgrind`
- Multi-dimensional arrays

# The Heap – Part II

# Refresher: memory organization

- Remember the way memory is organised?



# malloc example

```
#define NUM_ARRAYS 300
#define ARRAYSIZE 100000000 //100 million ints - 0.4 gb
void processArray(int * array)
{
    int j=0, sum=0;
    for (j=0; j < ARRAYSIZE; j++)
        sum += array[j];
    printf("%d\n",sum);
}
main()
{
    int* store_numbers;
    int i,j;
    for (i=0; i < NUM_ARRAYS;i++)
    {
        store_numbers = malloc(ARRAYSIZE*sizeof(int));
        store_numbers[0] = i;
        processArray(store_numbers);
        free(store_numbers);
    }
}
```

File:e20.c

malloc is defined in  
“stdlib.h”

← free the  
memory...!

# Memory allocation functions

- `void* malloc(size)`
  - `void*` is a 'void' pointer
  - it is a pointer to memory, but we don't know what is stored there
  - e.g., `int* a=malloc(sizeof(int)*10);`
  - allocation may fail: then `NULL` pointer returned
- `void* realloc(old_pointer, newsize)`
  - `old_pointer` must point to a previously allocated area.
  - contents remain unchanged up `min(newsize, oldsize)`.
  - If expanded, the contents of the new part are undefined.
  - After `realloc`, **`old_pointer` may become invalid.**
- `void* calloc(number, size);`
  - Guaranteed to **zero initialise** the memory
- `void free(pointer);`
  - Deallocates the space allocated by `malloc()`, `calloc()` or `realloc()`.

# Memory allocation functions

- `void* malloc(size)`
  - `void*` is a 'void' pointer
  - it is a pointer to memory, but we don't know what is stored there
  - e.g., `int* a=malloc(sizeof(int)*10);`
  - allocation may fail: then `NULL` pointer returned
- `void* realloc(old_pointer, newsize)`
  - `old_pointer` must point to a previously allocated area.
  - contents remain unchanged up `min(newsize, oldsize)`.
  - If expanded, the contents of the new part are undefined.
  - After `realloc`, `old_pointer` may become invalid.
- `void* calloc(number, size)`
  - Guaranteed to **zero initialise** the pointers?
- `void free(pointer);`
  - Deallocates the space allocated by `malloc()`, `calloc()` or `realloc()`.



# Casting types

- Casting (as in Java) converts one type of variable to another type
- In C, most primitive types can be moved without being cast
  - Even to smaller types
  - This would usually cause a compile error in Java
  - It is usually considered good practice to explicitly cast type conversions

```
main()
{
    int intvar;
    char charvar;
    long long longvar = (1024*1024+1) ;
    longvar = longvar * 4096;
    charvar = (char) longvar;
    intvar = (int)longvar;

    printf("%d %d %lld \n",charvar,intvar,longvar);
}
```

File: casting.c

# Casting Example

```
double db = 7 / 5;
```

```
printf("%f \n", db);
```

- prints 1.000000
- (7/5 is treated as an integer calculation)

```
double db = (double) (7 / 5);
```

```
printf("%f \n", db);
```

- Is exactly the same
- (7/5) is converted to a double after the calculation

```
double db = (double) 7 / 5;
```

```
printf("%f \n", db);
```

- prints 1.400000
- as 7 is treated as a double, hence 7/5 is no longer an integer calculation

```
double db = (double) (7 / 5.0);
```

```
printf("%f \n", db);
```

- prints 1.400000
- as 5 is treated as floating point, hence 7/5.0 is not an integer calculation

# Casting `void*`

- `malloc` returns a `void*`...
- `void*` cannot be dereferenced!

- this will not work:

```
int main(void)
{
    void* p=NULL;
    p = malloc(2*sizeof(int));
    printf("%d", *p);
}
```

`void_ptr.c`

- To dereference it: change to a different type of pointer!
  - So, cast it!
  - `int* a = (int*) malloc(sizeof(int)*10);`
- Why no problem before?
  - `int* a = malloc(sizeof(int)*10);`

# Casting General Pointers

- Pointers that are not void may also be explicitly cast to another type

- e.g.,

```
float number = 1.0;
float* number_pointer = &number;
int* int_pointer = (int*) number_pointer;
*int_pointer += 30;
printf("%f \n", number);
```

File:pointer\_mixing.c

- This is a dangerous thing to do...

# Arrays, Pointers & Malloc

- Arrays and pointers are mostly interchangeable...
- An **array** is a single, preallocated chunk of contiguous elements (all of the same type), fixed in size and location.
- A **pointer** is a reference to any data element (of a particular type) anywhere. A pointer must be assigned to point to space allocated elsewhere, but it can be reassigned (**and the space, if derived from malloc, can be resized**) at any time. A pointer can point to an array, **and can simulate (along with malloc) a dynamically allocated array**.
  - a pointer to a block of memory assigned by malloc is frequently treated (and can be referenced using `[ ]`) exactly as if it were a true array.

# Arrays, Pointers & Malloc

- Arrays and pointers are mostly interchangeable...
- An **array** is a single, preallocated chunk of contiguous elements (all of the same type), fixed in size and location
- A **pointer** is a reference to any  
anywhere. A pointer must be a  
elsewhere, but it can be reassigned  
**from malloc, can be resized**) a  
array, and can simulate (along with malloc, a dynamically  
**allocated array**.
  - a pointer to a block of memory assigned by malloc is frequently treated (and can be referenced using `[]`) exactly as if it were a true array.

Possible due to equivalence between  
pointer arithmetic and array indexing

# Heap or Stack?

- Where should you store data?
  - How much data?
  - How long do you need it for?
- Small amounts of data, that are only needed temporarily, should go on the stack (as local variables).
  - They are faster to use (no malloc to call, no pointer to dereference).
  - In some architectures, they may also be in faster memory (such as a pre-defined data cache).
- If you need to keep the data around for a while, or if it is a large amount, it needs to go on the heap.
  - But it must be freed after use

# Heap: common errors

- Not checking for allocation failures:
  - Memory allocation is not guaranteed to succeed. If there's no check for successful allocation implemented, this usually leads to a crash of the program or the entire system.
- Memory leaks:
  - Failure to deallocate memory using free leads to buildup of memory that is non-reusable memory, which is no longer used by the program. This wastes memory resources and can lead to allocation failures when these resources are exhausted.
- Logical errors:
  - All allocations must follow the same pattern:
  - **allocation** using malloc, **usage** to store data, **deallocation** using free.
  - Failures to adhere to this pattern, such as memory usage after a call to free or before a call to malloc, calling free twice, etc., usually leads to a crash of the program



# Example – 1

- Problem 1057

Title	Reverse input
Description	Reverse the input of a sequence of integers
Input	an integer $n$ , followed by $n$ integers
Output	the $n$ integers in reverse order
Sample Input	3 1 2 3
Sample Output	3 2 1
Hint	$n$ may be very large (up to 28 bits), only allocate the memory actually required

# Example – 1

```
#include <stdio.h>
#include <stdlib.h>
main()
{

    int i, array_size;
    int* array;
    scanf("%d",&array_size); // read the input size
    array = malloc(sizeof(int) * array_size); // allocate space for it

    for (i =0; i < array_size; i++) // read all the input
        scanf("%d",&array[i]);

    for (i =0; i < array_size; i++) // reverse the order
        printf("%d ",array[array_size - 1 - i]);

    free(array);
}
```

File:reverse\_malloc.c

# Review Heap Memory

- Memory allocation
  - Small, local, variables should go on the stack
  - Otherwise, you need to allocate heap memory at runtime
  - Use `malloc(size)` to allocate memory
  - `malloc` returns a pointer, or `NULL` if it can't be allocated
    - may need to **cast**
  - don't forget to free the memory after you have finished with it

# Heap debugging: Valgrind

# Using valgrind

- Valgrind can check for memory leaks!
- Compile with debug symbols (-g) and turn off optimization (-O0) for accurate reporting:
  - `gcc -o mem_leak -g -O0 -Wall mem_leak.c`
  - `valgrind --tool=memcheck --leak-check=full --show-leak-kinds=all ./mem_leak`

(demo)

- If you don't like all that typing... create a **Makefile**!

(demo)

- More info:
  - memcheck: <http://valgrind.org/docs/manual/quick-start.html>
  - other tools: <http://valgrind.org/info/tools.html>

# Reference: Interpreting Valgrind

```
valgrind --tool=memcheck --leak-check=full --show-leak-kinds=all ./mem_leak
```

```
==10927== Memcheck, a memory error detector
```

```
==10927== Copyright (C) 2002-2013, and GNU GPL'd, by Julian Seward et al.
```

```
==10927== Command: ./mem_leak
```

## <OUTPUT OF PROGRAM>

```
==10927==
```

```
==10927== HEAP SUMMARY:
```

```
==10927==      in use at exit: 120 bytes in 3 blocks
```

```
==10927==    total heap usage: 3 allocs, 0 frees, 120 bytes allocated
```

```
==10927==
```

```
==10927== 120 bytes in 3 blocks are definitely lost in loss record 1 of 1
```

```
==10927==    at 0x4C2AD10: calloc (vg_replace_malloc.c:623)
```

```
==10927==    by 0x4005C1: main (mem_leak.c:25)
```

```
==10927==
```

```
==10927== LEAK SUMMARY:
```

```
==10927==      definitely lost: 120 bytes in 3 blocks
```

```
==10927==      indirectly lost: 0 bytes in 0 blocks
```

```
==10927==      possibly lost: 0 bytes in 0 blocks
```

```
==10927==      still reachable: 0 bytes in 0 blocks
```

```
==10927==      suppressed: 0 bytes in 0 blocks
```

```
==10927==
```

```
==10927== For counts of detected and suppressed errors, rerun with: -v
```

```
==10927== ERROR SUMMARY: 1 errors from 1 contexts (suppressed: 0 from 0)
```

this is bad news

this tells you  
exactly where the  
memory was  
allocated

# Reference: Makefile

- contents of “Makefile”

```
mem_leak_debug:
    gcc -o mem_leak -g -O0 -Wall mem_leak.c

mem_check: mem_leak_debug
    valgrind --tool=memcheck --leak-check=full --show-leak-kinds=all ./mem_leak
```

- In terminal  
\$ make mem\_check
- Many other things makefiles can do!
  - ask google about “makefile tutorial”

# Reference: Makefile

- contents of “Makefile”

```
mem_leak_debug:
gcc -o mem_leak -g -O0 -Wall mem_leak.c

mem_check: mem_leak_debug
valgrind --tool=memcheck --leak-check=full --show-leak-kinds=all ./mem_leak
```

these are commands

these NEED to be tabs

this a dependency

- In terminal  
\$ make mem\_check
- Many other things makefiles can do!
  - ask google about “makefile tutorial”



# Review

- Heap memory allocation
  - use malloc etc.
  - and don't forget to free!
- Heap debugging... valgrind!
  - (yes, we all forget to free ;-)
- You now know how to...
  - ...deal with variable size memory requirements
  - ...decide between heap and stack
  - ...trace memory leaks
  - ...avoid typing long commands at the command line